

The k-means Algorithm

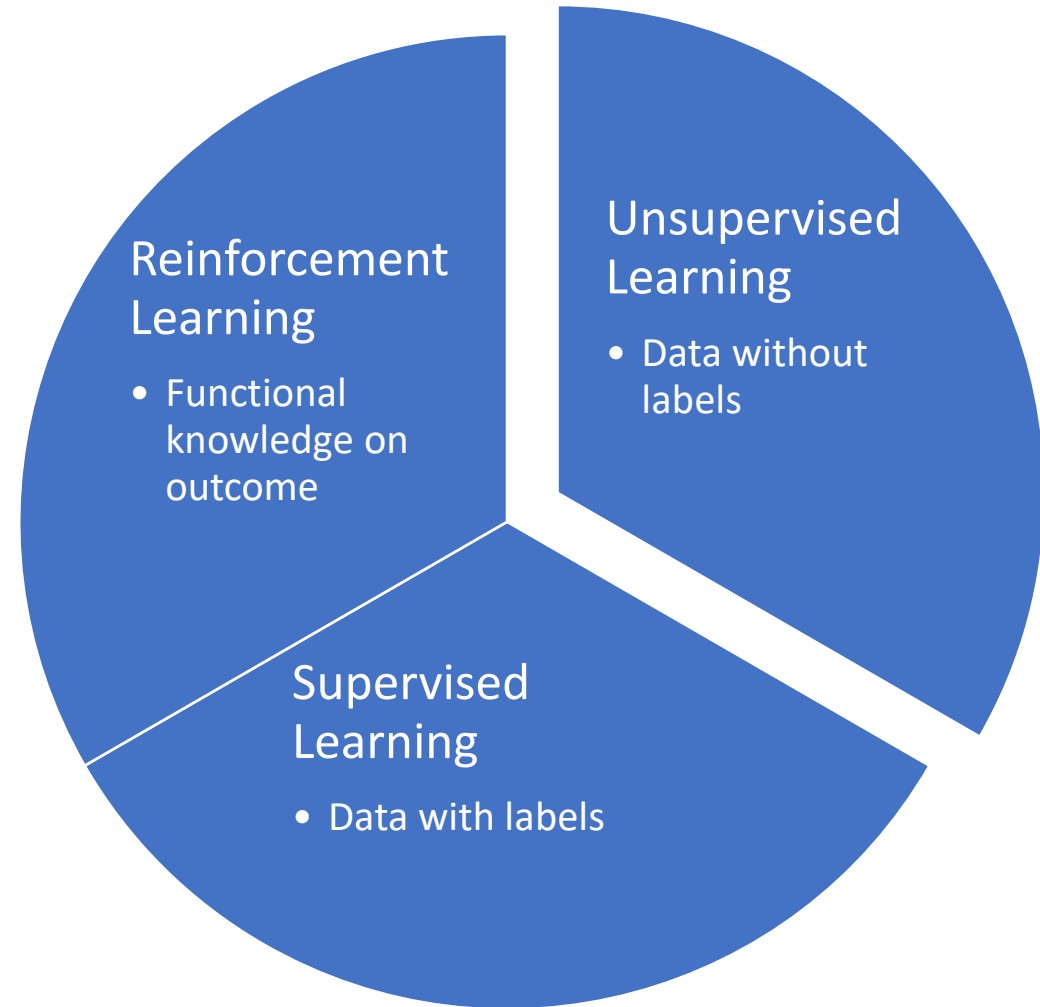
From data to insight

Overview


- Why k-means and what is it for?
- How does it work?
- Where to pay attention when using it?
- Where to look next?

K-means reveals structure

- Take a data set with numerical features as input
- No additional information (target, output) given.
- K-means shows where the data points „stick together“
 - These areas of data sticking together are called „clusters“



We start with data and many questions...



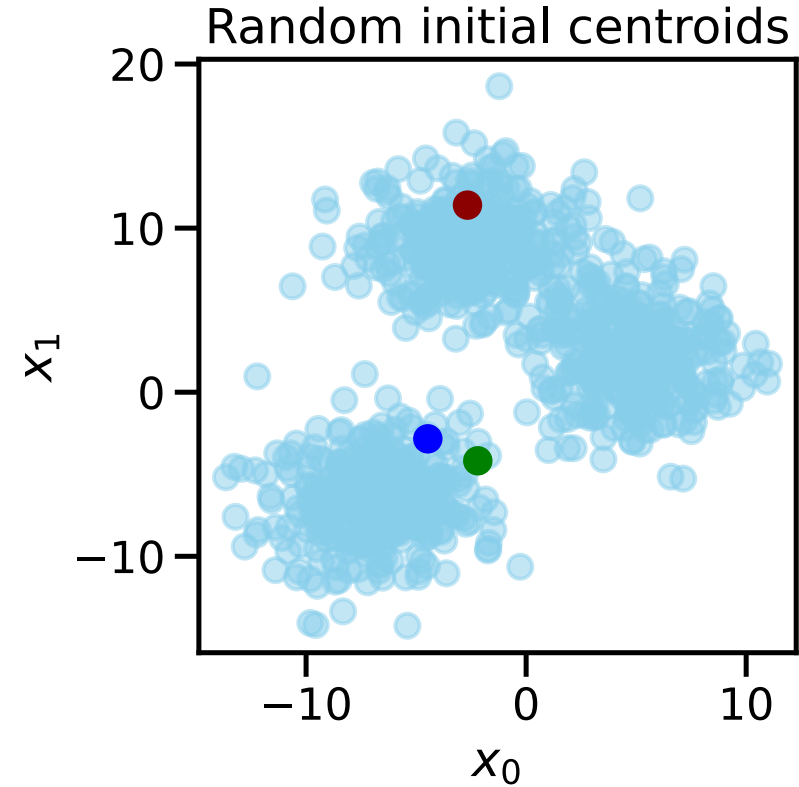
- We are given data in two dimensions:
 - This could be anything like temperature and humidity up to square meters and rent price.
- Without any additional information, what insight can we gather.
- Eyeballing in 2d reveals data lumps together...
- Generally: Number of features is the dimension of our space

The Algorithm (full)

```
centroids = initialise_centroids(points, k)
WHILE ((i < maxiter) & (delta_loss > eps)):
    forall point in points:
        forall centroid in centroids:
            calculate_distance(point, centroid)
        my_centroid = select_nearest_centroid(point)
        belonging_points[my_centroid].append(point)
    forall centroid in centroids:
        update_centroid(belonging_points[centroid])
    delta_loss = calculate_delta_loss(points, centroids)
    i++
```

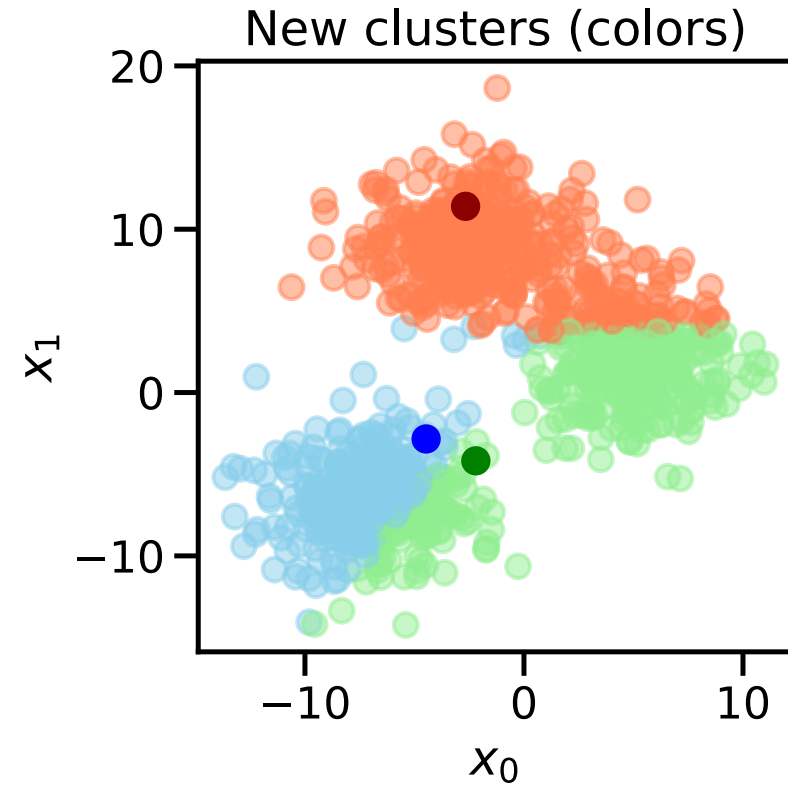
The Algorithm (initialise)

```
centroids = initialise_centroids(points, k)
```




The Algorithm (assign point to centroid)

```
centroids = initialise_centroids(points, k)
WHILE ((i < maxiter) & (delta_loss > eps)):
    forall point in points:
        forall centroid in centroids:
            calculate_distance(point, centroid)
        my_centroid = select_nearest_centroid(point)
        belonging_points[my_centroid].append(point)
```



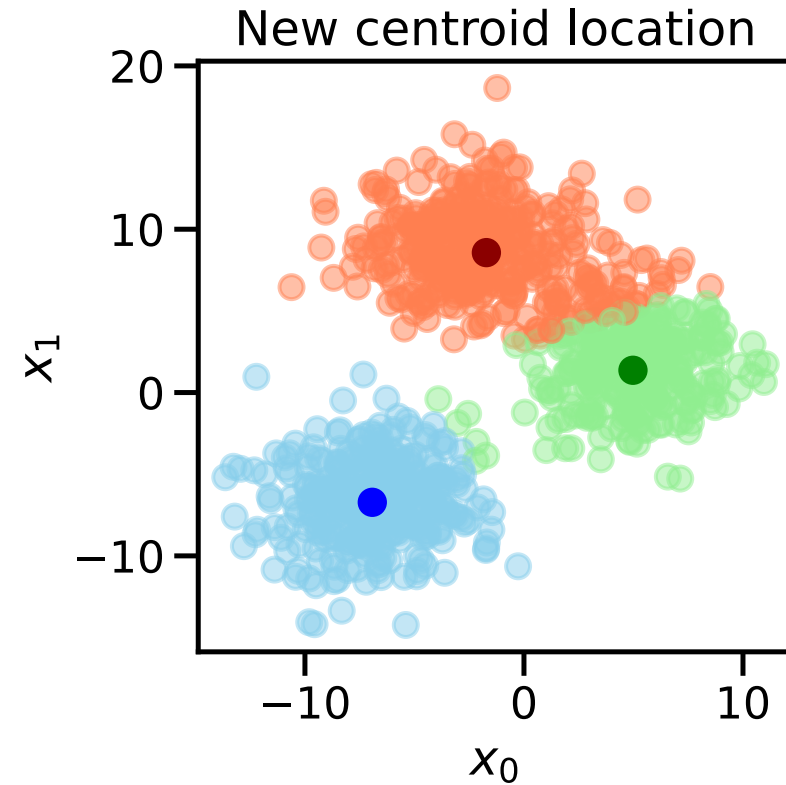
The Algorithm (assign point to centroid)

```
centroids = initialise_centroids(points, k)
WHILE ((i < maxiter) & (delta_loss > eps)):
    forall point in points:
        forall centroid in centroids:
            calculate_distance(point, centroid)
        my_centroid = select_nearest_centroid(point)
        belonging_points[my_centroid].append(point)
    forall centroid in centroids:
        update_centroid(belonging_points[centroid])
```

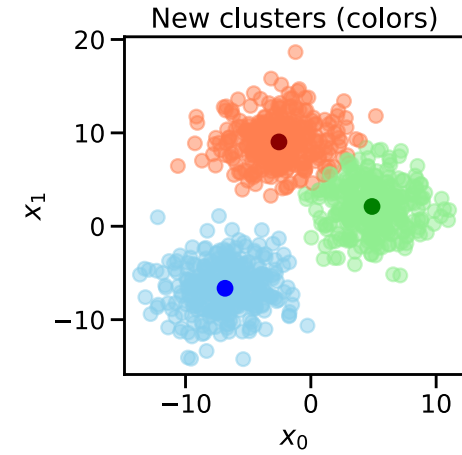





The Algorithm (prepare next iteration)

```
centroids = initialise_centroids(points, k)
WHILE ((i < maxiter) & (delta_loss > eps)):
    forall point in points:
        forall centroid in centroids:
            calculate_distance(point, centroid)
        my_centroid = select_nearest_centroid(point)
        belonging_points[my_centroid].append(point)
    forall centroid in centroids:
        update_centroid(belonging_points[centroid])
    delta_loss = calculate_delta_loss(points, centroids)
    i++
```



Next steps...



Have a look into the python example

```
i = 0
old_loss = np.inf
delta_loss = np.inf

while ((i < maxiter) & (delta_loss > eps)):

    # Calculate distance of all points to the all centroids
    for j, c in enumerate(centroids):
        distances[:, j] = get_distances(c, X)

    # Determine cluster membership of each point
    # by picking the closest centroid
    clusters = np.argmin(distances, axis=1)

    # Update centroid location using the newly
    # assigned data point cluster
    for c in range(k):
        centroids[c] = np.mean(X[clusters == c], 0)

    # For loss criterion calculate sum of squared distances to cluster centroid

    loss = calculate_loss(centroids, clusters)
    delta_loss = np.abs(old_loss - loss)
    old_loss = loss
    i = i+1
```

Calculations: distance and loss

- In classic k-means algorithm we use the euclidian distance
- For each point x and cluster with centroid C in n -dimensions:


$$d(x, C) = \sqrt{\sum_{i=1}^n (x_i - C_i)^2}$$

- The loss term is often called „inertia“ or „WCSS“ (within-cluster-sum-of-squares)
- It is the same (up to factor) as the sum over all clusters with centroids C_i of the variances within each cluster:



$$\sum_C \sum_{p \in C} d(p, C)^2$$

The one parameter: number of clusters k

- Has to be set before starting the loop.
- Can be derived from additional (business, science) insights.
- If totally unknown, there's the elbow-heuristic:
 - Try for consecutive k , save the loss after convergence
 - Plot losses depending on k
 - Pick one that is „the elbow“




Preparing your Hedge Fond

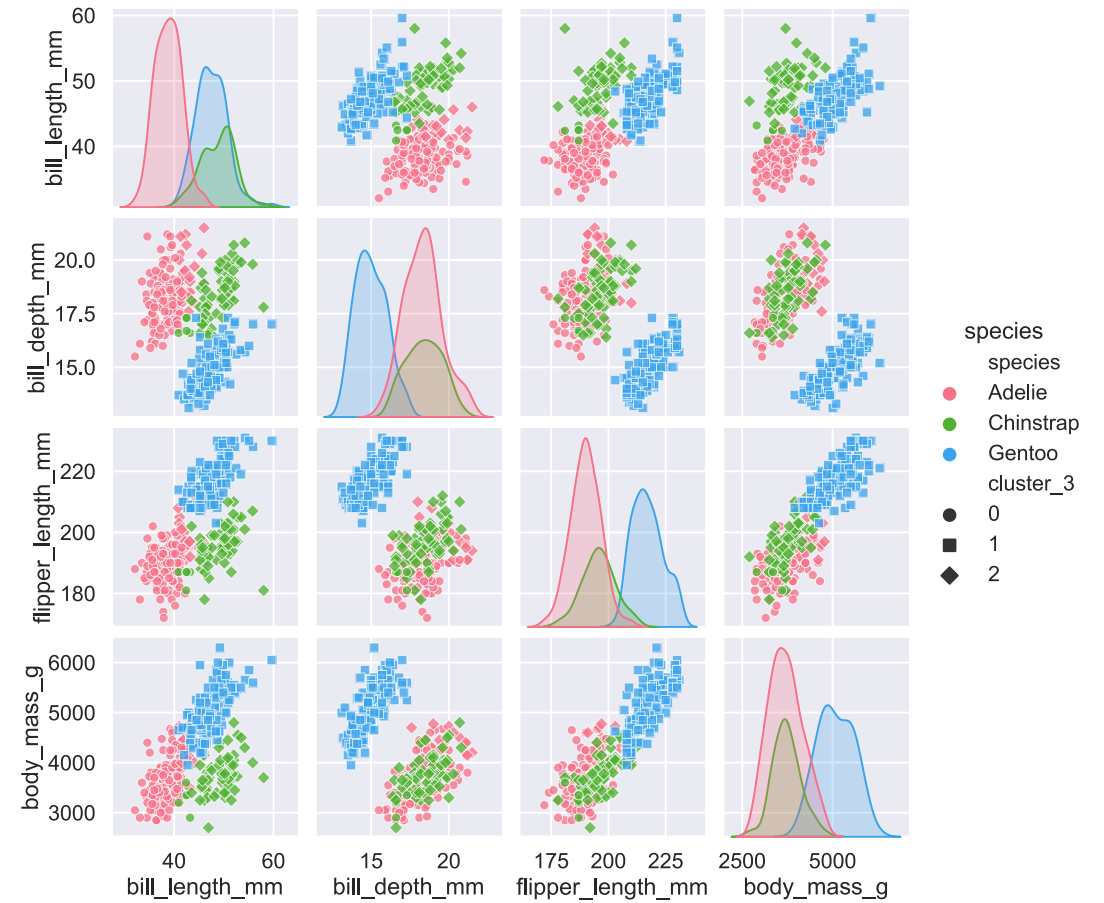



We choose $k=4$ based on our business Idea

- The elbow heuristic is useful, but if you have an idea how many clusters would be useful, go for it!
- Taking macro economic variables for regime prediction is often part of trading strategies...



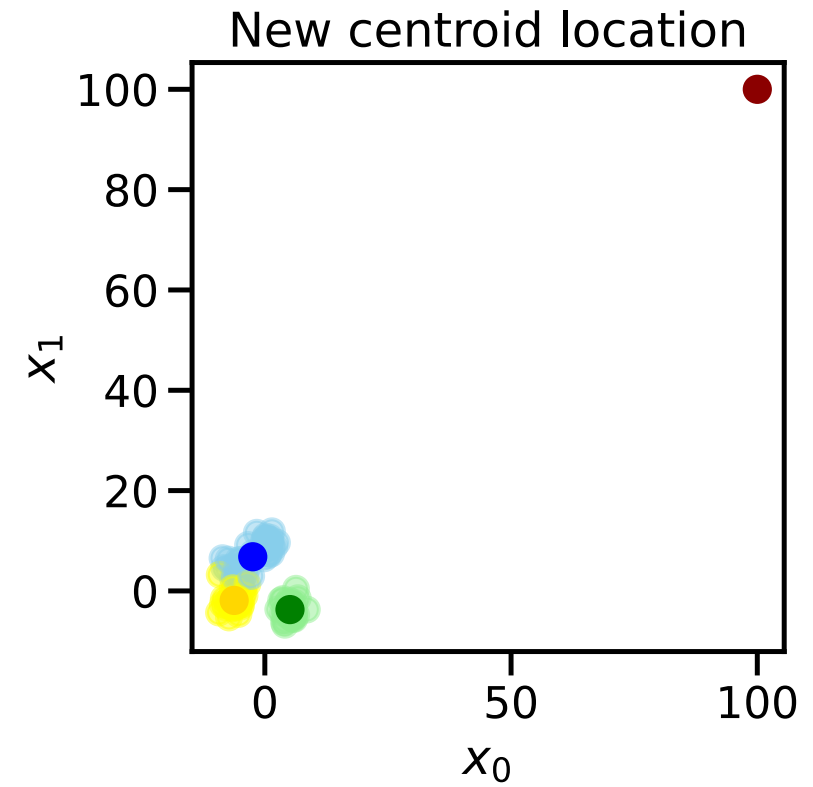
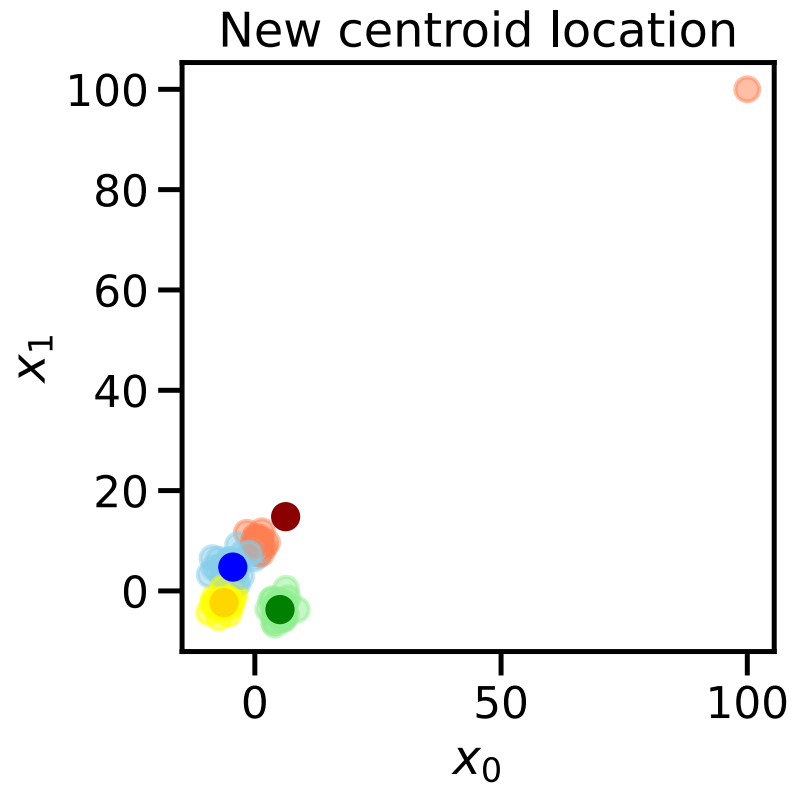
Re-discovering penguin species





Caveats:

- The algorithm works on features where calculating a mean makes sense – not on categorical data.
- The distance function ignores the scales of dimensions. Make sure the scales are what you want (it's a way, to weigh features!) or standardise.
- If you have severe outliers, the algorithm goes astray – Winsorisation (cutting off!) can help.
- Initializing the algorithm with random centroids makes it non deterministic – local minima can be compensated for by re-runs.
- The Algorithm works well, if data are in similar blobs.



Outliers



Not the same variance in blobs



Not even blobs...



Performance

- The k-means algorithm offers vast possibilities for parallelisation:
 - Using different initial clusters
 - Using different k (when finding a suitable one)
 - Distance of each point to each cluster
 - Calculation per cluster of new centroid
- Together with the simple math, its suited for GPU computing.

Rule of thumb: If calculations are somehow geometric and independent (game like explosions) they are suitable for GPUs.

Transparency

- The algorithm by itself is very transparent: we can see how it iterates to it's solution
- The math behind is well understood.
- The non-determinism needs to be carefully watched.
- In the selection of k and in the scaling can be hidden assumptions

Where to look next?

- Feel free to experiment with the python notebooks:
 - <https://github.com/jgerhard/k-means>
- If you want to try out real life financial data for your experiments:
 - <https://fred.stlouisfed.org/>
- On Kaggle you can find bunch of data sets and contests:
 - <https://kaggle.com/>
- If you want to play with basic data science algorithms, Joel Grus “Data Science from Scratch” can be a nice starter.
- If you prefer the next step, Jake VanderPlas’ “Python Data Science Handbook” is available freely: [Data Science Handbook](#)